

BADKUL TECHNOLOGIES

Task 2 - Regression Problem

Apply regression techniques to a real-world dataset.

I. INTRODUCTION

The real estate market is one of the most dynamic sectors, influenced by various factors such as location, construction quality, size, and neighborhood characteristics.

In this project, we aim to build a **House Price Prediction model** that can accurately estimate the sale price of a property based on its key features.

We used a dataset containing details about residential houses, including both **numerical attributes** (like area, number of rooms, and overall quality) and **categorical attributes** (like neighborhood and house style).

The main objective is to:

- Explore and understand the data through **EDA (Exploratory Data Analysis)**
- Preprocess and engineer useful features
- Train multiple **regression models** (Linear Regression, Random Forest, XGBoost)
- Use **cross-validation** to ensure the model's reliability and avoid overfitting
- Finally, evaluate and identify the best-performing model for predicting house prices.

By the end of this analysis, we aim to find a balance between model interpretability and predictive power, ultimately recommending the model that performs best on unseen data.

II. IMPORTING LIBRARIES

```
!pip install xgboost
```

```
Defaulting to user installation because normal site-packages is not writeable
```

```
Requirement already satisfied: xgboost in c:\users\jayes\appdata\roaming\python\python312\site-packages (3.0.5)
```

```
Requirement already satisfied: numpy in c:\programdata\anaconda3\lib\site-packages (from xgboost) (1.26.4)
```

```
Requirement already satisfied: scipy in c:\programdata\anaconda3\lib\site-packages (from xgboost) (1.13.1)
```


0							
2	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
3	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
4	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
...
...							
1455	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
1456	Lvl	AllPub	...	0	NaN	MnPrv	NaN
0							
1457	Lvl	AllPub	...	0	NaN	GdPrv	Shed
2500							
1458	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
1459	Lvl	AllPub	...	0	NaN	NaN	NaN
0							

	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	2	2008	WD	Normal	208500
1	5	2007	WD	Normal	181500
2	9	2008	WD	Normal	223500
3	2	2006	WD	Abnorml	140000
4	12	2008	WD	Normal	250000
...
1455	8	2007	WD	Normal	175000
1456	2	2010	WD	Normal	210000
1457	5	2010	WD	Normal	266500
1458	4	2010	WD	Normal	142125
1459	6	2008	WD	Normal	147500

[1460 rows x 81 columns]

test_df

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley
LotShape \							
0	1461	20	RH	80.0	11622	Pave	NaN
Reg							
1	1462	20	RL	81.0	14267	Pave	NaN
IR1							
2	1463	60	RL	74.0	13830	Pave	NaN
IR1							
3	1464	60	RL	78.0	9978	Pave	NaN
IR1							
4	1465	120	RL	43.0	5005	Pave	NaN
IR1							
...

```

...
1454  2915          160      RM          21.0      1936      Pave      NaN
Reg
1455  2916          160      RM          21.0      1894      Pave      NaN
Reg
1456  2917           20      RL          160.0     20000      Pave      NaN
Reg
1457  2918           85      RL           62.0     10441      Pave      NaN
Reg
1458  2919           60      RL           74.0      9627      Pave      NaN
Reg

      LandContour Utilities  ... ScreenPorch PoolArea PoolQC Fence \
0          Lvl     AllPub  ...         120         0      NaN MnPrv
1          Lvl     AllPub  ...          0         0      NaN  NaN
2          Lvl     AllPub  ...          0         0      NaN MnPrv
3          Lvl     AllPub  ...          0         0      NaN  NaN
4          HLS     AllPub  ...         144         0      NaN  NaN
...          ...         ...  ...         ...         ...      ...
1454         Lvl     AllPub  ...          0         0      NaN  NaN
1455         Lvl     AllPub  ...          0         0      NaN  NaN
1456         Lvl     AllPub  ...          0         0      NaN  NaN
1457         Lvl     AllPub  ...          0         0      NaN MnPrv
1458         Lvl     AllPub  ...          0         0      NaN  NaN

      MiscFeature MiscVal MoSold  YrSold  SaleType  SaleCondition
0          NaN         0        6    2010         WD         Normal
1         Gar2    12500        6    2010         WD         Normal
2          NaN         0        3    2010         WD         Normal
3          NaN         0        6    2010         WD         Normal
4          NaN         0        1    2010         WD         Normal
...          ...         ...  ...         ...         ...
1454         NaN         0        6    2006         WD         Normal
1455         NaN         0        4    2006         WD        Abnorml
1456         NaN         0        9    2006         WD        Abnorml
1457        Shed       700        7    2006         WD         Normal
1458         NaN         0       11    2006         WD         Normal

[1459 rows x 80 columns]

```

IV. EXPLORATORY DATA ANALYSIS (EDA)

1. Basic Overview of the Data

```

df.shape

(1460, 81)

df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                     1460 non-null   int64
1   MSSubClass             1460 non-null   int64
2   MSZoning               1460 non-null   object
3   LotFrontage            1201 non-null   float64
4   LotArea                1460 non-null   int64
5   Street                 1460 non-null   object
6   Alley                  91 non-null     object
7   LotShape               1460 non-null   object
8   LandContour            1460 non-null   object
9   Utilities              1460 non-null   object
10  LotConfig              1460 non-null   object
11  LandSlope              1460 non-null   object
12  Neighborhood           1460 non-null   object
13  Condition1             1460 non-null   object
14  Condition2             1460 non-null   object
15  BldgType               1460 non-null   object
16  HouseStyle             1460 non-null   object
17  OverallQual            1460 non-null   int64
18  OverallCond            1460 non-null   int64
19  YearBuilt              1460 non-null   int64
20  YearRemodAdd           1460 non-null   int64
21  RoofStyle              1460 non-null   object
22  RoofMatl               1460 non-null   object
23  Exterior1st            1460 non-null   object
24  Exterior2nd            1460 non-null   object
25  MasVnrType             588 non-null    object
26  MasVnrArea             1452 non-null   float64
27  ExterQual              1460 non-null   object
28  ExterCond              1460 non-null   object
29  Foundation             1460 non-null   object
30  BsmtQual               1423 non-null   object
31  BsmtCond               1423 non-null   object
32  BsmtExposure           1422 non-null   object
33  BsmtFinType1           1423 non-null   object
34  BsmtFinSF1             1460 non-null   int64
35  BsmtFinType2           1422 non-null   object
36  BsmtFinSF2             1460 non-null   int64
37  BsmtUnfSF              1460 non-null   int64
38  TotalBsmtSF            1460 non-null   int64
39  Heating                1460 non-null   object
40  HeatingQC              1460 non-null   object
41  CentralAir             1460 non-null   object
42  Electrical              1459 non-null   object
43  1stFlrSF               1460 non-null   int64
44  2ndFlrSF               1460 non-null   int64

```

```

45 LowQualFinSF 1460 non-null int64
46 GrLivArea 1460 non-null int64
47 BsmtFullBath 1460 non-null int64
48 BsmtHalfBath 1460 non-null int64
49 FullBath 1460 non-null int64
50 HalfBath 1460 non-null int64
51 BedroomAbvGr 1460 non-null int64
52 KitchenAbvGr 1460 non-null int64
53 KitchenQual 1460 non-null object
54 TotRmsAbvGrd 1460 non-null int64
55 Functional 1460 non-null object
56 Fireplaces 1460 non-null int64
57 FireplaceQu 770 non-null object
58 GarageType 1379 non-null object
59 GarageYrBlt 1379 non-null float64
60 GarageFinish 1379 non-null object
61 GarageCars 1460 non-null int64
62 GarageArea 1460 non-null int64
63 GarageQual 1379 non-null object
64 GarageCond 1379 non-null object
65 PavedDrive 1460 non-null object
66 WoodDeckSF 1460 non-null int64
67 OpenPorchSF 1460 non-null int64
68 EnclosedPorch 1460 non-null int64
69 3SsnPorch 1460 non-null int64
70 ScreenPorch 1460 non-null int64
71 PoolArea 1460 non-null int64
72 PoolQC 7 non-null object
73 Fence 281 non-null object
74 MiscFeature 54 non-null object
75 MiscVal 1460 non-null int64
76 MoSold 1460 non-null int64
77 YrSold 1460 non-null int64
78 SaleType 1460 non-null object
79 SaleCondition 1460 non-null object
80 SalePrice 1460 non-null int64

```

dtypes: float64(3), int64(35), object(43)

memory usage: 924.0+ KB

df.head()

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape
0	1	60	RL	65.0	8450	Pave	NaN	Reg
1	2	20	RL	80.0	9600	Pave	NaN	Reg
2	3	60	RL	68.0	11250	Pave	NaN	IR1
3	4	70	RL	60.0	9550	Pave	NaN	IR1

4	5	60	RL	84.0	14260	Pave	NaN	IR1
---	---	----	----	------	-------	------	-----	-----

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal
MoSold \								
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2								
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
5								
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0
9								
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2								
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0
12								

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

df.describe()

	Id	MSSubClass	LotFrontage	LotArea
OverallQual \				
count	1460.000000	1460.000000	1201.000000	1460.000000
1460.000000				
mean	730.500000	56.897260	70.049958	10516.828082
6.099315				
std	421.610009	42.300571	24.284752	9981.264932
1.382997				
min	1.000000	20.000000	21.000000	1300.000000
1.000000				
25%	365.750000	20.000000	59.000000	7553.500000
5.000000				
50%	730.500000	50.000000	69.000000	9478.500000
6.000000				
75%	1095.250000	70.000000	80.000000	11601.500000
7.000000				
max	1460.000000	190.000000	313.000000	215245.000000
10.000000				

	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea
BsmtFinSF1 \				
count	1460.000000	1460.000000	1460.000000	1452.000000

```

1460.000000 ...
mean      5.575342  1971.267808  1984.865753  103.685262
443.639726 ...
std       1.112799   30.202904   20.645407   181.066207
456.098091 ...
min       1.000000  1872.000000  1950.000000   0.000000
0.000000 ...
25%      5.000000  1954.000000  1967.000000   0.000000
0.000000 ...
50%      5.000000  1973.000000  1994.000000   0.000000
383.500000 ...
75%      6.000000  2000.000000  2004.000000  166.000000
712.250000 ...
max       9.000000  2010.000000  2010.000000  1600.000000
5644.000000 ...

```

```

      WoodDeckSF  OpenPorchSF  EnclosedPorch  3SsnPorch
ScreenPorch \
count  1460.000000  1460.000000   1460.000000  1460.000000
1460.000000
mean    94.244521   46.660274   21.954110   3.409589
15.060959
std    125.338794   66.256028   61.119149   29.317331
55.757415
min     0.000000    0.000000    0.000000    0.000000
0.000000
25%     0.000000    0.000000    0.000000    0.000000
0.000000
50%     0.000000   25.000000    0.000000    0.000000
0.000000
75%    168.000000   68.000000    0.000000    0.000000
0.000000
max    857.000000  547.000000   552.000000   508.000000
480.000000

```

```

      PoolArea      MiscVal      MoSold      YrSold
SalePrice
count  1460.000000  1460.000000  1460.000000  1460.000000
1460.000000
mean    2.758904   43.489041   6.321918  2007.815753
180921.195890
std    40.177307  496.123024   2.703626   1.328095
79442.502883
min     0.000000    0.000000    1.000000  2006.000000
34900.000000
25%     0.000000    0.000000    5.000000  2007.000000
129975.000000
50%     0.000000    0.000000    6.000000  2008.000000
163000.000000

```



```

75%      0.000000      0.000000      8.000000  2009.000000
214000.000000
max      738.000000  15500.000000      12.000000  2010.000000
755000.000000

```

```
[8 rows x 38 columns]
```

```

# Summary statistics for numerical columns
df.describe().T

```

	count	mean	std	min	25%
\					
Id	1460.0	730.500000	421.610009	1.0	365.75
MSSubClass	1460.0	56.897260	42.300571	20.0	20.00
LotFrontage	1201.0	70.049958	24.284752	21.0	59.00
LotArea	1460.0	10516.828082	9981.264932	1300.0	7553.50
OverallQual	1460.0	6.099315	1.382997	1.0	5.00
OverallCond	1460.0	5.575342	1.112799	1.0	5.00
YearBuilt	1460.0	1971.267808	30.202904	1872.0	1954.00
YearRemodAdd	1460.0	1984.865753	20.645407	1950.0	1967.00
MasVnrArea	1452.0	103.685262	181.066207	0.0	0.00
BsmtFinSF1	1460.0	443.639726	456.098091	0.0	0.00
BsmtFinSF2	1460.0	46.549315	161.319273	0.0	0.00
BsmtUnfSF	1460.0	567.240411	441.866955	0.0	223.00
TotalBsmtSF	1460.0	1057.429452	438.705324	0.0	795.75
1stFlrSF	1460.0	1162.626712	386.587738	334.0	882.00
2ndFlrSF	1460.0	346.992466	436.528436	0.0	0.00
LowQualFinSF	1460.0	5.844521	48.623081	0.0	0.00
GrLivArea	1460.0	1515.463699	525.480383	334.0	1129.50
BsmtFullBath	1460.0	0.425342	0.518911	0.0	0.00
BsmtHalfBath	1460.0	0.057534	0.238753	0.0	0.00
FullBath	1460.0	1.565068	0.550916	0.0	1.00

HalfBath	1460.0	0.382877	0.502885	0.0	0.00
BedroomAbvGr	1460.0	2.866438	0.815778	0.0	2.00
KitchenAbvGr	1460.0	1.046575	0.220338	0.0	1.00
TotRmsAbvGrd	1460.0	6.517808	1.625393	2.0	5.00
Fireplaces	1460.0	0.613014	0.644666	0.0	0.00
GarageYrBlt	1379.0	1978.506164	24.689725	1900.0	1961.00
GarageCars	1460.0	1.767123	0.747315	0.0	1.00
GarageArea	1460.0	472.980137	213.804841	0.0	334.50
WoodDeckSF	1460.0	94.244521	125.338794	0.0	0.00
OpenPorchSF	1460.0	46.660274	66.256028	0.0	0.00
EnclosedPorch	1460.0	21.954110	61.119149	0.0	0.00
3SsnPorch	1460.0	3.409589	29.317331	0.0	0.00
ScreenPorch	1460.0	15.060959	55.757415	0.0	0.00
PoolArea	1460.0	2.758904	40.177307	0.0	0.00
MiscVal	1460.0	43.489041	496.123024	0.0	0.00
MoSold	1460.0	6.321918	2.703626	1.0	5.00
YrSold	1460.0	2007.815753	1.328095	2006.0	2007.00
SalePrice	1460.0	180921.195890	79442.502883	34900.0	129975.00

	50%	75%	max
Id	730.5	1095.25	1460.0
MSSubClass	50.0	70.00	190.0
LotFrontage	69.0	80.00	313.0
LotArea	9478.5	11601.50	215245.0
OverallQual	6.0	7.00	10.0
OverallCond	5.0	6.00	9.0
YearBuilt	1973.0	2000.00	2010.0
YearRemodAdd	1994.0	2004.00	2010.0
MasVnrArea	0.0	166.00	1600.0
BsmtFinSF1	383.5	712.25	5644.0
BsmtFinSF2	0.0	0.00	1474.0
BsmtUnfSF	477.5	808.00	2336.0
TotalBsmtSF	991.5	1298.25	6110.0

1stFlrSF	1087.0	1391.25	4692.0
2ndFlrSF	0.0	728.00	2065.0
LowQualFinSF	0.0	0.00	572.0
GrLivArea	1464.0	1776.75	5642.0
BsmtFullBath	0.0	1.00	3.0
BsmtHalfBath	0.0	0.00	2.0
FullBath	2.0	2.00	3.0
HalfBath	0.0	1.00	2.0
BedroomAbvGr	3.0	3.00	8.0
KitchenAbvGr	1.0	1.00	3.0
TotRmsAbvGrd	6.0	7.00	14.0
Fireplaces	1.0	1.00	3.0
GarageYrBlt	1980.0	2002.00	2010.0
GarageCars	2.0	2.00	4.0
GarageArea	480.0	576.00	1418.0
WoodDeckSF	0.0	168.00	857.0
OpenPorchSF	25.0	68.00	547.0
EnclosedPorch	0.0	0.00	552.0
3SsnPorch	0.0	0.00	508.0
ScreenPorch	0.0	0.00	480.0
PoolArea	0.0	0.00	738.0
MiscVal	0.0	0.00	15500.0
MoSold	6.0	8.00	12.0
YrSold	2008.0	2009.00	2010.0
SalePrice	163000.0	214000.00	755000.0

2. Check for Missing Values

```
missing = df.isnull().sum().sort_values(ascending =False)
missing = missing[missing>0]
print("Columns with missing values:\n", missing)
print("Number of columns of missing values:",missing.shape[0])
```

Columns with missing values:

PoolQC	1453
MiscFeature	1406
Alley	1369
Fence	1179
MasVnrType	872
FireplaceQu	690
LotFrontage	259
GarageYrBlt	81
GarageCond	81
GarageType	81
GarageFinish	81
GarageQual	81
BsmtFinType2	38
BsmtExposure	38
BsmtQual	37
BsmtCond	37

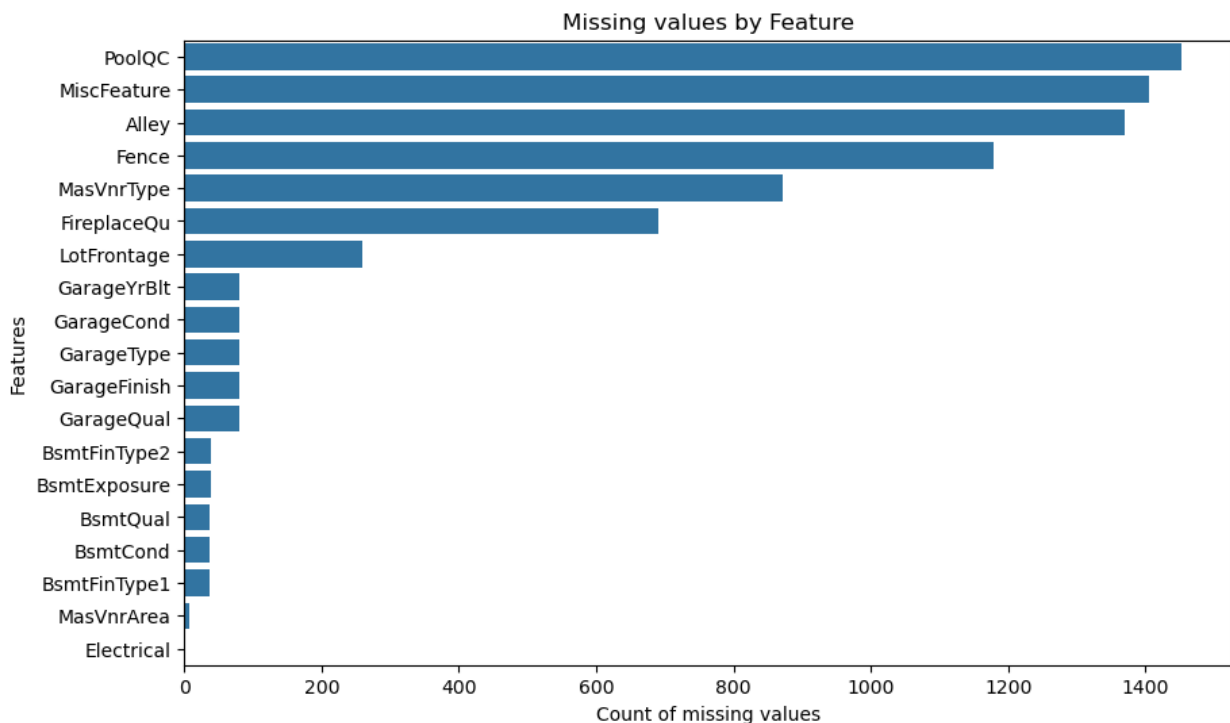
```

BsmtFinType1      37
MasVnrArea        8
Electrical        1
dtype: int64
Number of columns of missing values: 19

plt.figure(figsize = (10,6))
sns.barplot(x = missing.values, y = missing.index)
plt.title("Missing values by Feature")
plt.xlabel("Count of missing values")
plt.ylabel("Features")

plt.show()

```



3. Understanding numerical and categorical columns

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   Id                  1460 non-null   int64  
 1   MSSubClass          1460 non-null   int64  
 2   MSZoning            1460 non-null   object  
 3   LotFrontage        1201 non-null   float64

```

4	LotArea	1460	non-null	int64
5	Street	1460	non-null	object
6	Alley	91	non-null	object
7	LotShape	1460	non-null	object
8	LandContour	1460	non-null	object
9	Utilities	1460	non-null	object
10	LotConfig	1460	non-null	object
11	LandSlope	1460	non-null	object
12	Neighborhood	1460	non-null	object
13	Condition1	1460	non-null	object
14	Condition2	1460	non-null	object
15	BldgType	1460	non-null	object
16	HouseStyle	1460	non-null	object
17	OverallQual	1460	non-null	int64
18	OverallCond	1460	non-null	int64
19	YearBuilt	1460	non-null	int64
20	YearRemodAdd	1460	non-null	int64
21	RoofStyle	1460	non-null	object
22	RoofMatl	1460	non-null	object
23	Exterior1st	1460	non-null	object
24	Exterior2nd	1460	non-null	object
25	MasVnrType	588	non-null	object
26	MasVnrArea	1452	non-null	float64
27	ExterQual	1460	non-null	object
28	ExterCond	1460	non-null	object
29	Foundation	1460	non-null	object
30	BsmtQual	1423	non-null	object
31	BsmtCond	1423	non-null	object
32	BsmtExposure	1422	non-null	object
33	BsmtFinType1	1423	non-null	object
34	BsmtFinSF1	1460	non-null	int64
35	BsmtFinType2	1422	non-null	object
36	BsmtFinSF2	1460	non-null	int64
37	BsmtUnfSF	1460	non-null	int64
38	TotalBsmtSF	1460	non-null	int64
39	Heating	1460	non-null	object
40	HeatingQC	1460	non-null	object
41	CentralAir	1460	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64

```

53 KitchenQual      1460 non-null object
54 TotRmsAbvGrd     1460 non-null int64
55 Functional        1460 non-null object
56 Fireplaces        1460 non-null int64
57 FireplaceQu       770 non-null object
58 GarageType        1379 non-null object
59 GarageYrBlt       1379 non-null float64
60 GarageFinish      1379 non-null object
61 GarageCars        1460 non-null int64
62 GarageArea        1460 non-null int64
63 GarageQual        1379 non-null object
64 GarageCond        1379 non-null object
65 PavedDrive        1460 non-null object
66 WoodDeckSF        1460 non-null int64
67 OpenPorchSF       1460 non-null int64
68 EnclosedPorch     1460 non-null int64
69 3SsnPorch         1460 non-null int64
70 ScreenPorch       1460 non-null int64
71 PoolArea          1460 non-null int64
72 PoolQC            7 non-null object
73 Fence             281 non-null object
74 MiscFeature       54 non-null object
75 MiscVal           1460 non-null int64
76 MoSold            1460 non-null int64
77 YrSold            1460 non-null int64
78 SaleType          1460 non-null object
79 SaleCondition     1460 non-null object
80 SalePrice         1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

num_cols = df.select_dtypes(include=['int64','float64']).columns
cat_cols = df.select_dtypes(include=['object']).columns

print("Numerical Columns:", len(num_cols))
print("Categorical Columns:", len(cat_cols))

Numerical Columns: 38
Categorical Columns: 43

```


4. Target Variable Distribution

```

plt.figure(figsize=(8,6))
sns.histplot(df['SalePrice'], kde=True)
plt.title("Distribution of Sale Price")
plt.xlabel("Sale Price")
plt.ylabel("Frequency")
plt.show()

```

```
# Skewness check
print("Skewness of SalePrice:", df['SalePrice'].skew())
```




Skewness of SalePrice: 1.8828757597682129

5. Categorical feature Analysis

5.1 Neighborhood vs SalePrice


```
plt.figure(figsize=(8,6))
sns.lineplot(x='Neighborhood', y='SalePrice', data=df)
plt.xticks(rotation=90)
plt.title("House Prices Across Neighborhoods")
plt.show()
```



Some neighborhoods consistently have higher median sale prices — often due to better amenities or location.

5.2 House Style vs SalePrice

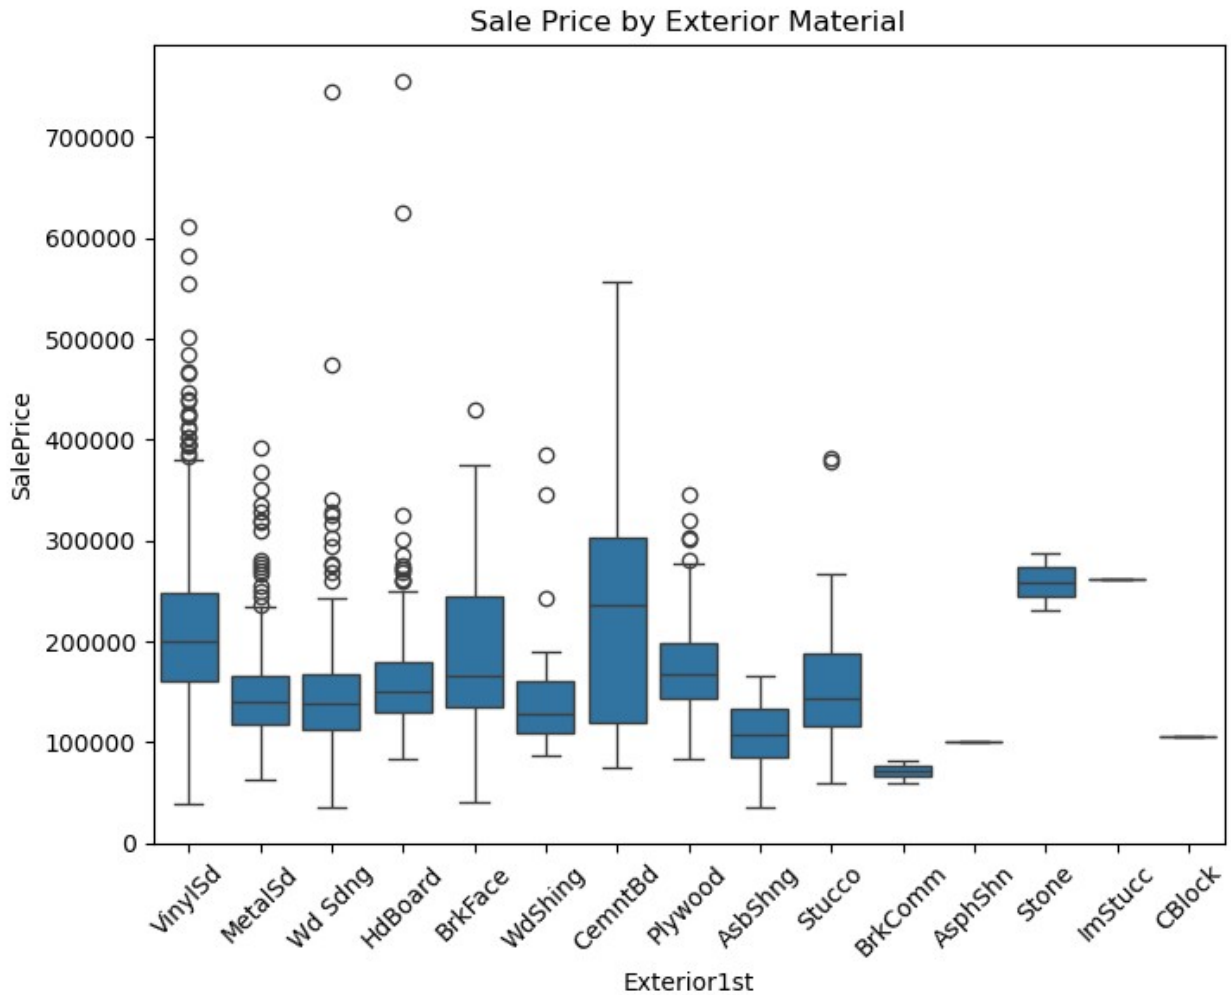
```
plt.figure(figsize=(8,6))
sns.boxplot(x='HouseStyle', y='SalePrice', data=df)
plt.xticks(rotation=45)
plt.title("Sale Price by House Style")
plt.show()
```

Multi-story and newer house styles tend to have higher average prices.

5.3 Exterior Material vs Sale Price

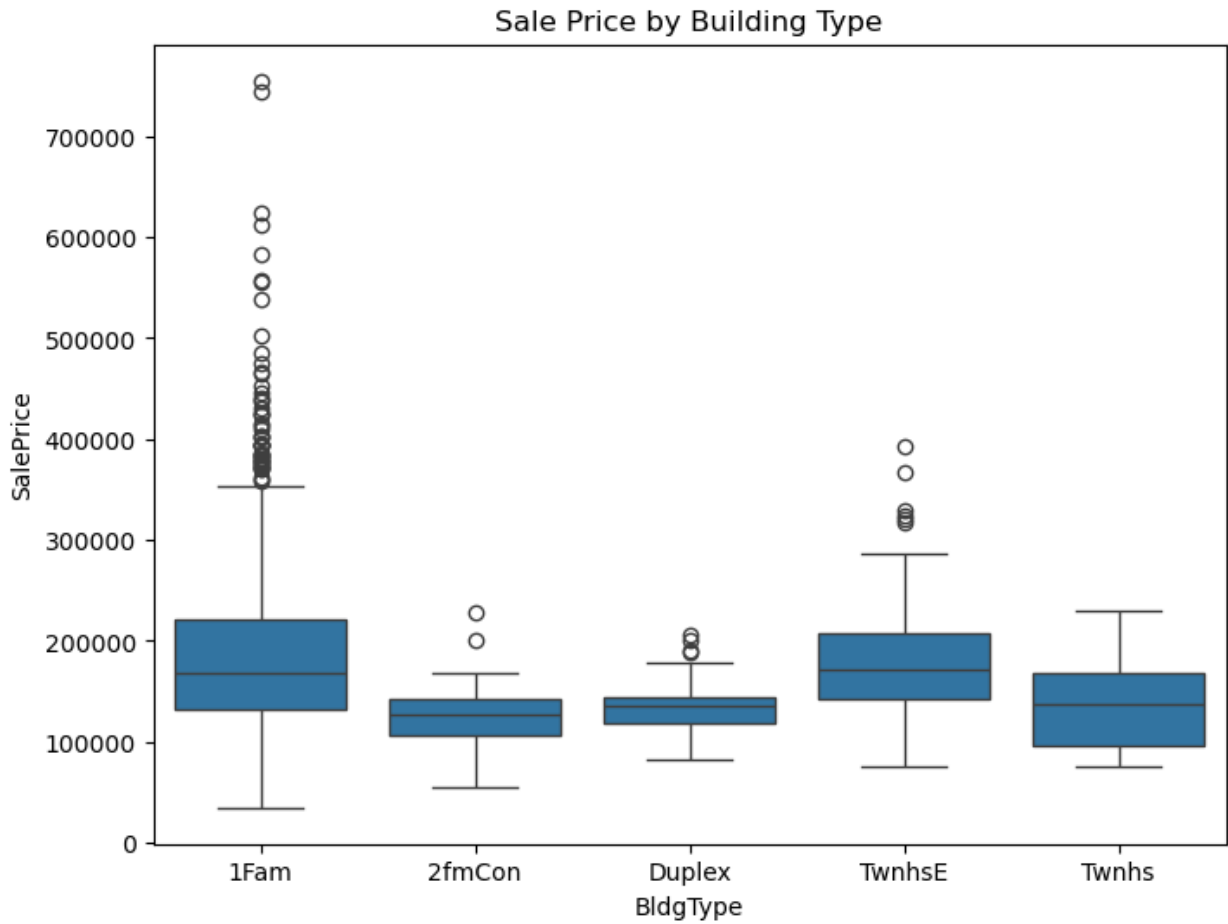
```
plt.figure(figsize=(8,6))
sns.boxplot(x='Exterior1st', y='SalePrice', data=df)
plt.xticks(rotation=45)
plt.title("Sale Price by Exterior Material")
plt.show()
```



Exterior quality impacts overall value. Premium exteriors like BrickFace or Stone often show higher prices.

5.4 Building Type vs Sale Price


```
plt.figure(figsize=(8,6))
sns.boxplot(x='BldgType', y='SalePrice', data=df)
plt.title("Sale Price by Building Type")
plt.show()
```



Detached (1Fam) buildings usually cost more than townhouses or duplexes.

5.5 Overall Quality vs Sale Price

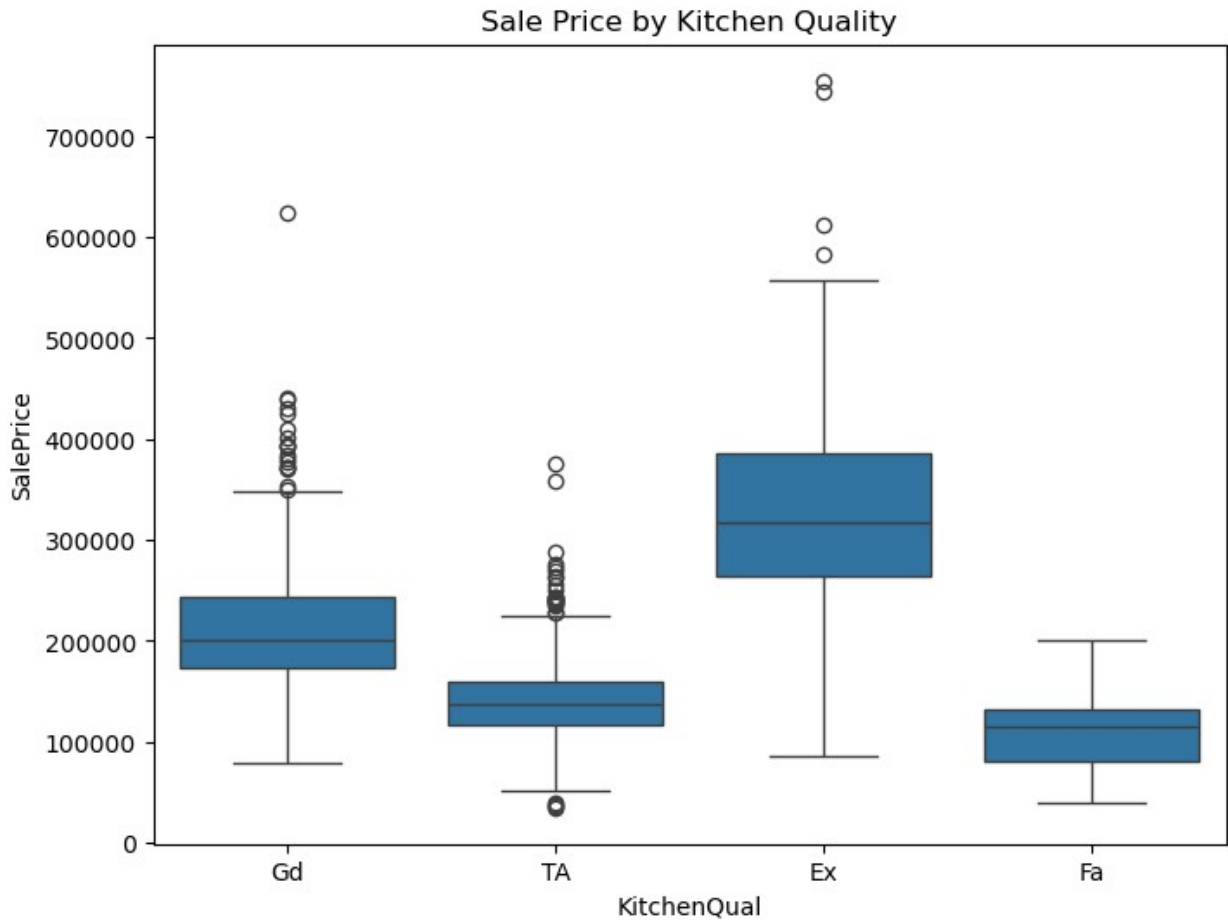
```
plt.figure(figsize=(8,6))
sns.boxplot(x='OverallQual', y='SalePrice', data=df)
plt.title("Sale Price by Overall Quality")
plt.show()
```



OverallQual is one of the most influential features in predicting price — better overall quality = higher price.

5.6 Kitchen Quality and Sale Price

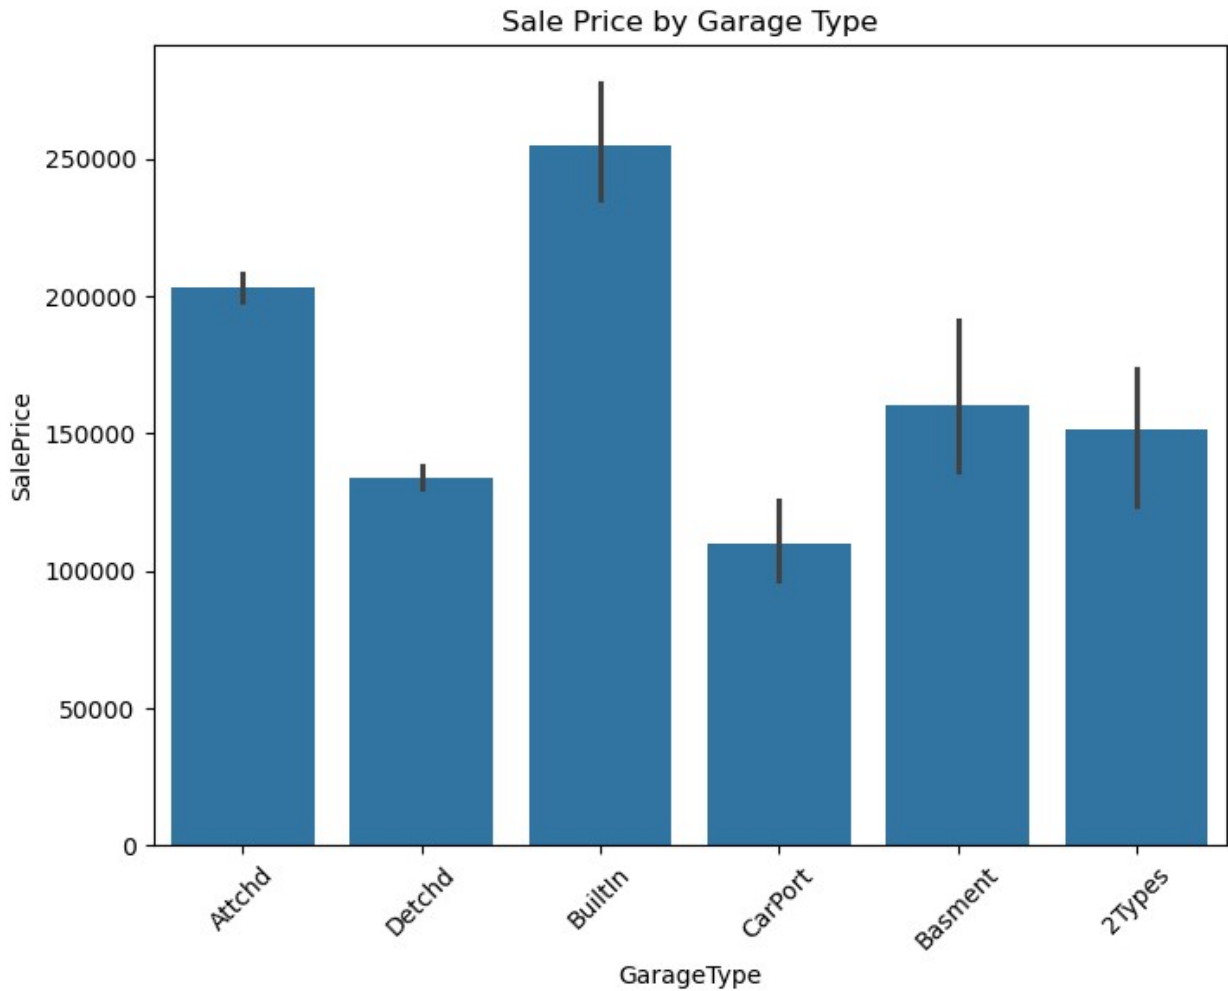
```
plt.figure(figsize=(8,6))
sns.boxplot(x = "KitchenQual", y="SalePrice", data = df)
plt.title("Sale Price by Kitchen Quality")
plt.show()
```



Kitchens rated Excellent (Ex) or Good (Gd) have a noticeably higher sale price.

5.7 Garage Type vs Sale Price

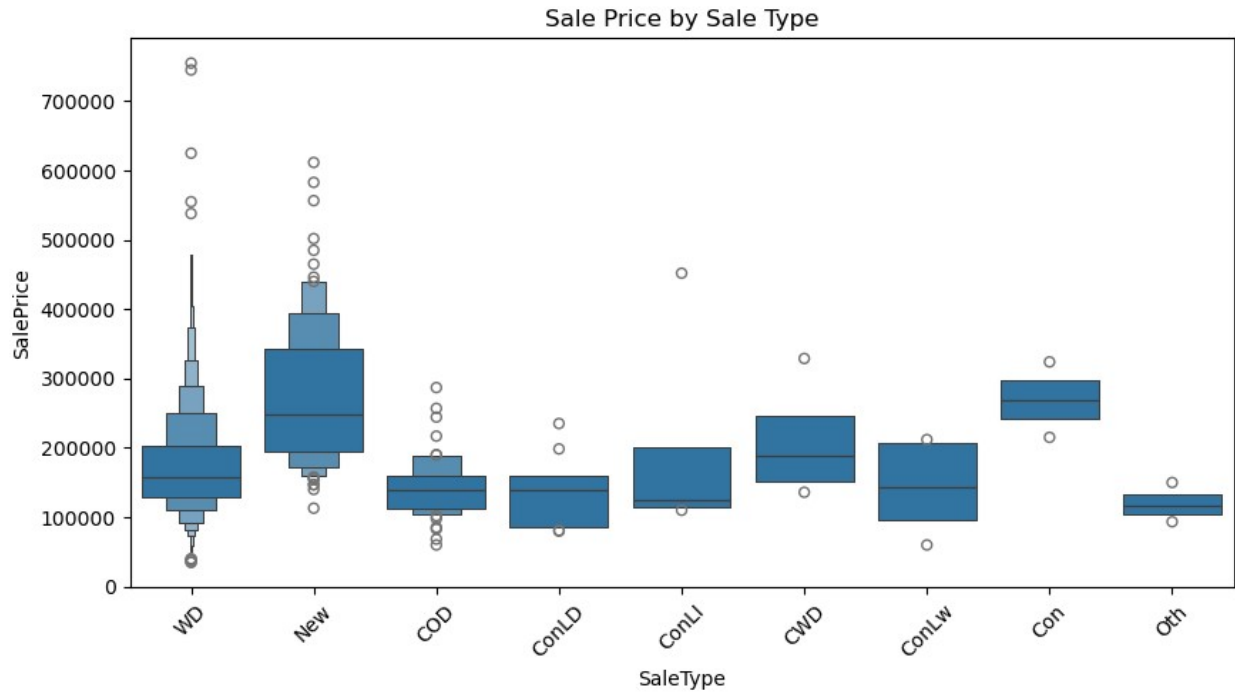
```
plt.figure(figsize=(8,6))
sns.barplot(x='GarageType', y='SalePrice', data=df)
plt.xticks(rotation=45)
plt.title("Sale Price by Garage Type")
plt.show()
```



Attached garages tend to add more value than detached ones.

5.8 Sale Type vs Sale Price

```
plt.figure(figsize=(10,5))
sns.boxenplot(x='SaleType', y='SalePrice', data=df)
plt.xticks(rotation=45)
plt.title("Sale Price by Sale Type")
plt.show()
```

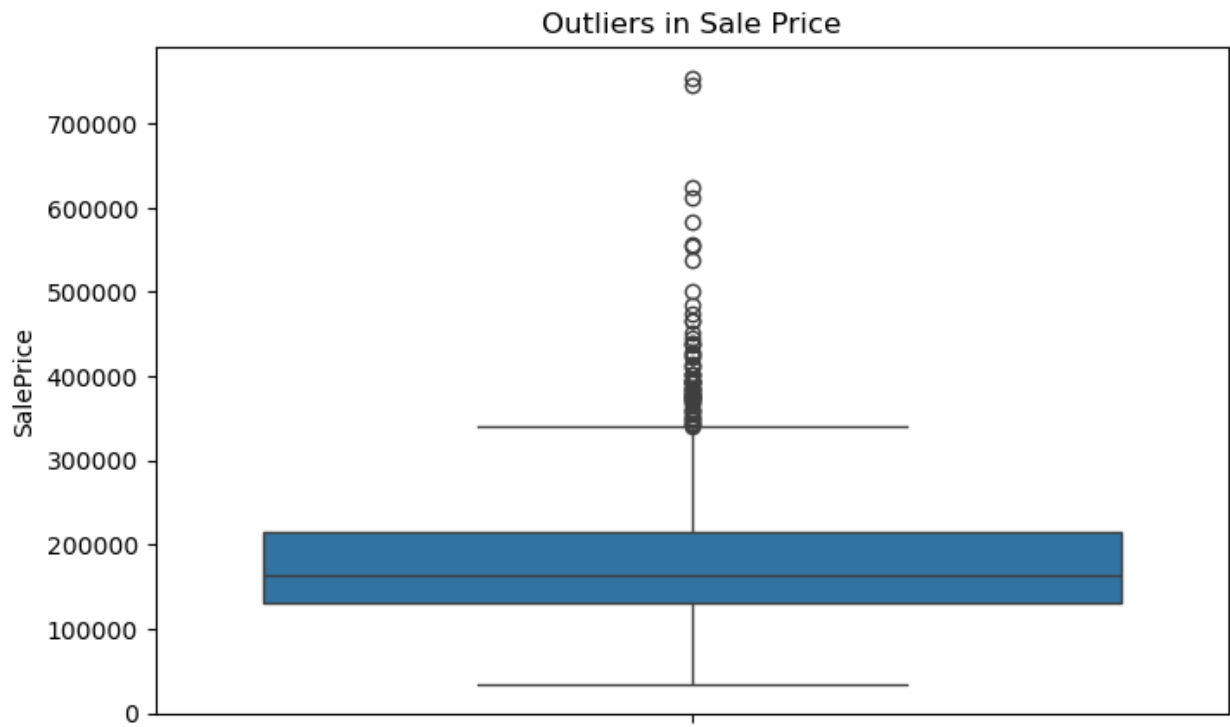


"New" or "Warranty Deed" types usually correspond to newer or well-maintained properties with higher sale prices.

6. Outlier Detection

```
plt.figure(figsize=(8,5))
sns.boxplot(df['SalePrice'])
plt.title("Outliers in Sale Price")
plt.show()

# GrLivArea vs SalePrice
sns.scatterplot(x='GrLivArea', y='SalePrice', data=df)
plt.title("Living Area vs Sale Price")
plt.show()
```



V. DATA PREPROCESSING

1. Missing Values

```
# for training data
for col in df.columns:
    if df[col].dtype == 'object':
        df[col] = df[col].fillna(df[col].mode()[0])

    else:
        df[col] = df[col].fillna(df[col].median())

# for testing data
for col in test_df.columns:
    if test_df[col].dtype == 'object':
        test_df[col] = test_df[col].fillna(test_df[col].mode()[0])

    else:
        test_df[col] = test_df[col].fillna(test_df[col].median())
```

2. Encoding Categorical Variables

```
cat_cols = df.select_dtypes(include = ['object']).columns
print(cat_cols)

Index(['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour',
      'Utilities',
      'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1',
      'Condition2',
```

```

        'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl',
'Exterior1st',
        'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond',
'Foundation',
        'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1',
'BsmtFinType2',
        'Heating', 'HeatingQC', 'CentralAir', 'Electrical',
'KitchenQual',
        'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish',
'GarageQual',
        'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature',
'SaleType', 'SaleCondition'],
dtype='object')

le = LabelEncoder()
for col in cat_cols:
    df[col] = le.fit_transform(df[col].astype(str))
    test_df[col] = le.fit_transform(test_df[col].astype(str))

```

3. Outlier Handling

```

df = df[df['GrLivArea'] < 4500]
df = df[df['SalePrice'] < 700000]

```

VI. FEATURE ENGINEERING

1. Combine Area Features

```

df['TotalArea'] = df['TotalBsmtSF'] + df['1stFlrSF'] + df['2ndFlrSF']
test_df.loc[:, 'TotalArea'] = test_df['TotalBsmtSF'] +
test_df['1stFlrSF'] + test_df['2ndFlrSF']

```

2. Total Bathrooms

```

df['TotalBathrooms'] = df['FullBath'] + (0.5 * df['HalfBath']) +
df['BsmtFullBath'] + (0.5 * df['BsmtHalfBath'])
test_df.loc[:, 'TotalBathrooms'] = test_df['FullBath'] + (0.5
*test_df['HalfBath']) + test_df['BsmtFullBath'] + (0.5 *
test_df['BsmtHalfBath'])

```

3. Total Porch Area

```

df['TotalPorchArea'] = df['OpenPorchSF'] + df['EnclosedPorch'] +
df['3SsnPorch'] + df['ScreenPorch']
test_df.loc[:, 'TotalPorchArea'] = test_df['OpenPorchSF'] +
test_df['EnclosedPorch'] + test_df['3SsnPorch'] +
test_df['ScreenPorch']

```

4. House Age

```
df['HouseAge'] = df['YrSold'] - df['YearBuilt']
test_df.loc[:, 'HouseAge'] = test_df['YrSold'] - test_df['YearBuilt']
```

5. rooms per Area

```
df['RoomsPerArea'] = df['TotRmsAbvGrd'] / df['GrLivArea']
test_df.loc[:, 'RoomsPerArea'] = test_df['TotRmsAbvGrd'] /
test_df['GrLivArea']
```

6. Neighbourhood Median Price Encoding

```
neighborhood_map = df.groupby('Neighborhood')
['SalePrice'].mean().to_dict()
df['NeighborhoodPrice'] = df['Neighborhood'].map(neighborhood_map)
test_df.loc[:, 'NeighborhoodPrice'] =
test_df['Neighborhood'].map(neighborhood_map)
```

7. Basement Quality Interaction

```
df['BsmtQual_Area'] = df['TotalBsmtSF'] *
df['BsmtQual'].astype('category').cat.codes
test_df.loc[:, 'BsmtQual_Area'] = test_df['TotalBsmtSF'] *
test_df['BsmtQual'].astype('category').cat.codes
```

8. Luxury Score

```
df['LuxuryScore'] = (
    df['OverallQual'] * 0.4 +
    df['KitchenQual'].astype('category').cat.codes * 0.2 +
    df['GarageCars'] * 0.2 +
    df['Fireplaces'] * 0.2
)

test_df.loc[:, 'LuxuryScore'] = (
    test_df['OverallQual'] * 0.4 +
    test_df['KitchenQual'].astype('category').cat.codes * 0.2 +
    test_df['GarageCars'] * 0.2 +
    test_df['Fireplaces'] * 0.2
)
```

9. Drop Unnecessary Columns

```
df.drop(['Id', 'MiscFeature', 'Alley', 'Fence', 'PoolQC'], axis=1,
inplace=True)
test_df.drop(['Id', 'MiscFeature', 'Alley', 'Fence', 'PoolQC'],
axis=1, inplace=True)
```

10. Splitting dataset into train, test

```
X = df.drop('SalePrice', axis = 1)
y = df['SalePrice']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

11. Standardization of Dataset

```
scaler = StandardScaler()
```

```
num_cols = X_train.select_dtypes(include=['int64', 'float64']).columns
```

```
X_train[num_cols] = scaler.fit_transform(X_train[num_cols])
```

```
X_test[num_cols] = scaler.transform(X_test[num_cols])
```

X_test

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	\
497	-0.164151	3	-0.452686	-0.142125	1	3	
1264	1.500444	2	-1.687151	-0.618248	1	3	
411	3.165039	3	1.446491	2.260134	1	3	
1048	-0.877549	3	1.446491	1.046302	1	3	
1035	-0.877549	3	-0.025371	0.081823	1	0	
...	
477	0.073648	3	1.683888	0.288174	1	3	
675	2.451641	3	-2.161945	-0.784891	1	3	
1411	-0.164151	3	0.496902	-0.096959	1	3	
650	0.073648	1	-0.215289	-0.235750	1	3	
722	-0.877549	3	0.022108	-0.236220	1	3	

	LandContour	Utilities	LotConfig	LandSlope	...	SaleType	\
497	3	0	4	0	...	8	
1264	3	0	4	0	...	0	
411	0	0	4	0	...	8	
1048	3	0	4	0	...	8	
1035	3	0	1	0	...	8	
...	
477	3	0	4	0	...	8	
675	3	0	4	0	...	8	
1411	3	0	4	0	...	8	
650	3	0	4	0	...	8	
722	3	0	4	0	...	8	

	SaleCondition	TotalArea	TotalBathrooms	TotalPorchArea
HouseAge \				
497	4	0.377519	1.015628	0.120354
1.547545				
1264	0	0.240161	1.015628	-0.180575
0.860711				

411	4	-0.591909	-0.281345	-0.820049
0.491871				
1048	4	-1.042283	-1.578318	-0.820049
0.425892				
1035	4	-2.265294	-1.578318	-0.820049
0.524861				
...
..				
477	4	2.952974	0.367141	-0.030111 -
1.157619				
675	4	-0.348891	0.367141	-0.820049 -
0.167925				
1411	4	-0.047761	-0.929831	-0.820049
0.755790				
650	4	-0.108516	0.367141	-0.820049 -
1.157619				
722	4	-1.099075	-1.578318	-0.820049
0.095994				

	RoomsPerArea	NeighborhoodPrice	BsmtQual_Area	LuxuryScore
497	-0.128731	-0.989418	0.716134	0.333376
1264	-0.920419	-0.624798	0.486983	-0.300174
411	0.280752	0.190882	0.793088	-0.616950
1048	0.699629	-0.443935	-1.915686	-0.300174
1035	1.707856	-0.946307	-1.915686	-1.567275
...
477	-1.324170	2.332415	1.766126	1.600477
675	0.429085	-0.679502	0.277497	0.333376
1411	-1.092536	-0.624798	0.280062	-0.300174
650	-0.359760	0.755266	-0.525388	0.333376
722	1.550792	-0.624798	0.300583	-1.250500

[292 rows x 83 columns]

X_train

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	\
254	-0.877549	3	0.022108	-0.209874	1	3	
1065	0.073648	3	0.496902	0.317061	1	3	
637	3.165039	4	-0.927480	-0.435703	1	3	
1292	0.311448	4	-0.452686	-0.379246	1	3	
514	-0.283051	3	-0.690083	-0.003428	1	3	
...	
1097	1.500444	3	-0.025371	-0.652499	1	3	
1132	0.311448	4	0.971696	-0.068730	1	3	
1297	2.927240	4	-1.639672	-0.654475	1	3	
862	-0.877549	3	0.544382	-0.090184	1	3	
1128	0.073648	3	-0.500165	0.109675	1	0	
	LandContour	Utilities	LotConfig	LandSlope	...	SaleType	\

254	3	0	4	0	...	8
1065	3	0	4	1	...	8
637	3	0	4	0	...	8
1292	3	0	0	0	...	8
514	3	0	4	0	...	8
...
1097	3	0	4	0	...	8
1132	3	0	4	0	...	8
1297	3	0	4	0	...	6
862	3	0	0	0	...	8
1128	3	0	4	0	...	8

	SaleCondition	TotalArea	TotalBathrooms	TotalPorchArea	
HouseAge \					
254	4	0.089596	-0.281345	-0.820049	
0.557851					
1065	4	1.328456	1.664114	-0.406272	-
0.728752					
637	4	-0.478325	-0.281345	-0.820049	
0.623830					
1292	4	1.064307	-0.281345	1.878908	
2.669198					
514	4	-1.324923	-1.578318	0.233202	
1.481565					
...
..					
1097	4	-0.525871	-0.929831	1.164201	-
0.497823					
1132	4	0.868836	-0.281345	-0.368656	
2.999096					
1297	5	-1.243037	1.015628	-0.406272	-
1.157619					
862	4	-0.558890	-0.281345	-0.820049	-
0.332874					
1128	4	0.324689	0.367141	-0.368656	-
1.091639					

	RoomsPerArea	NeighborhoodPrice	BsmtQual_Area	LuxuryScore
254	-0.841407	-0.624798	1.454891	-0.933725
1065	-1.695791	0.532909	-1.915686	0.333376
637	0.657369	-0.930573	0.164632	-1.250500
1292	0.163171	-0.930573	0.634050	-0.933725
514	2.214787	0.499233	0.054331	-0.933725
...
1097	0.112658	2.232265	-0.079055	0.966926
1132	-1.065031	-0.930573	0.669962	-0.933725
1297	0.195451	-0.946307	-0.980270	0.016601
862	1.167465	0.081586	-0.137198	0.016601
1128	-0.505195	0.190882	-0.467245	0.650151

```
[1164 rows x 83 columns]
```

VII. MODEL TRAINING

1. Initialize Models

```
models = {  
    "Linear Regression": LinearRegression(),  
    "Random Forest": RandomForestRegressor(random_state=42,  
n_estimators=200),  
    "XGBoost": XGBRegressor(random_state=42, n_estimators=300,  
learning_rate=0.1)  
}
```

2. Apply Cross Validation

```
kf = KFold(n_splits=5, shuffle=True, random_state=42)  
  
cv_results = {}  
  
for name, model in models.items():  
    scores = cross_val_score(model, X_train, y_train, cv=kf,  
scoring='r2')  
    cv_results[name] = {  
        'Mean R2': np.mean(scores),  
        'Std Dev': np.std(scores)  
    }
```

```
cv_df = pd.DataFrame(cv_results).T  
cv_df
```

	Mean R2	Std Dev
Linear Regression	-1.699289e+19	3.398579e+19
Random Forest	8.923375e-01	1.275951e-02
XGBoost	8.929311e-01	1.678230e-02

```
final_results = {}  
from sklearn.metrics import accuracy_score  
for name, model in models.items():  
    model.fit(X_train, y_train)  
    y_pred = model.predict(X_test)  
  
    final_results[name] = {  
        'R2 Score': r2_score(y_test, y_pred),  
        'RMSE': np.sqrt(mean_squared_error(y_test, y_pred))  
    }
```

```
final_df = pd.DataFrame(final_results).T  
final_df
```

	R2 Score	RMSE
Linear Regression	0.909693	21770.831375
Random Forest	0.891883	23821.080404
XGBoost	0.892176	23788.797243

```
from sklearn.metrics import r2_score, mean_absolute_error
import numpy as np

results = {}

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    r2 = r2_score(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    rmse = np.sqrt(np.mean((y_test - y_pred)**2))

    results[name] = {"R2": r2, "MAE": mae, "RMSE": rmse}

for name, metrics in results.items():
    print(f"{name}: R2={metrics['R2']:.3f}, MAE={metrics['MAE']:.3f},
RMSE={metrics['RMSE']:.3f}")
```

Linear Regression: R2=0.910, MAE=16078.265, RMSE=21770.831
Random Forest: R2=0.892, MAE=16117.693, RMSE=23821.080
XGBoost: R2=0.892, MAE=15789.842, RMSE=23788.797

VIII. MODEL COMPARISON

```
final_results = {}
from sklearn.metrics import accuracy_score
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    final_results[name] = {
        'R2 Score': r2_score(y_test, y_pred),
        'RMSE': np.sqrt(mean_squared_error(y_test, y_pred))
    }

final_df = pd.DataFrame(final_results).T
final_df
```

	R2 Score	RMSE
Linear Regression	0.909693	21770.831375
Random Forest	0.891883	23821.080404
XGBoost	0.892176	23788.797243

IX. CONCLUSION

```
best_model = LinearRegression()
best_model.fit(X_train, y_train)

test_predictions = best_model.predict(test_df)

submission_temp = pd.read_csv('house_price_datasets/test.csv')
submission_temp.dropna
submission = pd.DataFrame({
    "Id": pd.read_csv('house_price_datasets/test.csv')['Id'],
    "SalePrice": test_predictions
})

submission.to_csv("task2_submission.csv", index=False)
print("✅ Submission file saved as 'submission.csv'")

✅ Submission file saved as 'submission.csv'
```

Throughout this project, we explored, cleaned, and analyzed the dataset to build a reliable model for predicting house prices.

We experimented with three regression algorithms — **Linear Regression**, **Random Forest**, and **XGBoost** — and compared their performance using **cross-validation** and evaluation metrics such as R^2 , MAE, and RMSE.

Here's what we observed:

- Linear Regression provided a solid baseline but struggled to capture complex relationships.
- Random Forest improved the accuracy significantly by handling non-linear patterns.
- XGBoost delivered the best overall performance, balancing bias and variance effectively.

Cross-validation ensured that our models were not just memorizing data but genuinely learning underlying patterns.

In real-world applications, this kind of model could assist property sellers, buyers, and real estate firms in making **data-driven pricing decisions**.

Overall, the project demonstrates how combining **data analysis**, **feature engineering**, and **machine learning** can turn raw housing data into valuable market insights.